THOMAS YEOH

+1 (669)-799-7334 | tyeoh@berkeley.edu | https://tyeoh9.github.io/tyeoh/

Education

University of California, Irvine Expected: Dec 2026
Master of Computer Science

University of California, Berkeley Graduated: May 2025
B.S. Electrical Engineering and Computer Science

- Coursework: Data Structures & Algorithms, Graphics, Machine Learning, Databases, Computer Vision, Computer Architecture, Probability & Optimization
- Activities & Awards: Cal Leadership Scholarship, Google Developers Club (VP of Strategy)

Skills

Languages: Python, Java, C++, HTML, CSS, JavaScript, SQL, R, Bash, Assembly, Scheme
Technologies: Linux, GCP, PyTorch, Pandas, NumPy, Git, Flask, React, Node.js, Streamlit, Docker, Jupyter

Experiences

Zoofi AI Sunnyvale, CA
Software Engineer Intern June 2025 – Aug 2025

- Designed and built the core Python backend infrastructure and agentic platform using Google ADK to deliver the MVP and attract the first 3 beta users.
- Automated document parsing and extraction using Google ADK for client onboarding, cutting manual effort by 60%.
- Implemented FastAPI endpoints with Pydantic models for validation, ensuring a consistent and robust API schema.

Berkeley Operations & Behavioral Analytics Lab Berkeley, CA
Research Assistant Jan 2024 – May 2025

- Performed data analysis and manipulation using Pandas on behavioral datasets to validate empirical findings.
- Optimized mathematical models and the data pipeline, accelerating pre-computation speeds by 44% using efficient algorithms such as dynamic programming.
- Leveraged HPC clusters in Unix/Linux environments for parallel and distributed computing.

Keysight Technologies Penang, Malaysia
Software Engineer Intern May 2024 – August 2024

- Designed and shipped a web app, leveraging NLP to analyze request forms and reduce human review time by 25%.
- Built an end-to-end data pipeline with Python, Flask, and data structures like vectorized arrays to efficiently process over 120,000 data points daily.
- Utilized various libraries (Keras, Sklearn) for advanced data modeling and analysis on large datasets.
- Experimented with ML methods: XGBoost, SVM, K-means, logistic regression, and Facebook fastText.

Projects

RookieDB (Java, SQL)

- Implemented a SQLite-like database in Java using policies such as multi-granularity locking and ARIES recovery.
- Utilized data structures like B+ trees and join algorithms, improving information retrieval and query execution speeds by ~30%.

Non-photorealistic Shaders Engine (C++, GLSL)

- Built a C++ and OpenGL non-photorealistic renderer using object-oriented design, featuring multiple shaders.
- Integrated ImGui for runtime parameterization of shaders, enabling flexible and user-friendly rendering control.